

✓ Muhammad Tahir K214503

LAB 04

✓ 1. Initially perform the following three exercises on the following Links:

<https://www.analyticsvidhya.com/blog/2021/01/a-guide-to-the-naive-bayes-algorithm/>

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import sklearn
```

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import sklearn

dataset = pd.read_csv('Social_Network_Ads.csv')
X = dataset.iloc[:, [1, 2, 3]].values
y = dataset.iloc[:, -1].values

print(X)
```

```
[[ 'Male' 19 19000]
 [ 'Male' 35 20000]
 [ 'Female' 26 43000]
 ...
 [ 'Female' 50 20000]
 [ 'Male' 36 33000]
 [ 'Female' 49 36000]]
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
X[:,0] = le.fit_transform(X[:,0])
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state = 0)
```

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(X_train, y_train)
```

```

GaussianNB
GaussianNB()
```

```
y_pred = classifier.predict(X_test)
```

```
y_pred
```

```
array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1,
       0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
       1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1,
       0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1])
```

y_test

```
array([[0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
        0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
        1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1,
        0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1]])
```

```
from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test, y_pred)
ac = accuracy_score(y_test, y_pred)
```

cm

```
array([[56,  2],
       [ 4, 18]])
```

ac

```
0.925
```

<https://www.statology.org/k-fold-cross-validation-in-python/>

```
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LinearRegression
from numpy import mean
from numpy import absolute
from numpy import sqrt
import pandas as pd
```

```
df = pd.DataFrame({'y': [6, 8, 12, 14, 14, 15, 17, 22, 24, 23],
                   'x1': [2, 5, 4, 3, 4, 6, 7, 5, 8, 9],
                   'x2': [14, 12, 12, 13, 7, 8, 7, 4, 6, 5]})
```

```
#define predictor and response variables
```

```
X = df[['x1', 'x2']]
y = df['y']
```

```
#define cross-validation method to use
```

```
cv = KFold(n_splits=10, random_state=1, shuffle=True)
```

```
#build multiple linear regression model
```

```
model = LinearRegression()
```

```
#use k-fold CV to evaluate model
```

```
scores = cross_val_score(model, X, y, scoring='neg_mean_absolute_error',
                          cv=cv, n_jobs=-1)
```

```
#view mean absolute error
```

```
mean(absolute(scores))
```

```
3.146154808346972
```

```
#define predictor and response variables
```

```
X = df[['x1', 'x2']]
y = df['y']
```

```
#define cross-validation method to use
```

```
cv = KFold(n_splits=5, random_state=1, shuffle=True)
```

```
#build multiple linear regression model
```

```
model = LinearRegression()
```

```
#use LOOCV to evaluate model
scores = cross_val_score(model, X, y, scoring='neg_mean_squared_error',
                          cv=cv, n_jobs=-1)

#view RMSE
sqrt(mean(absolute(scores)))
```

→ 4.284373111711817

<https://www.datacamp.com/community/tutorials/naive-bayes-scikit-learn>

```
# Assigning features and label variables
Weather = ['Sunny', 'Sunny', 'Overcast', 'Rainy', 'Rainy', 'Rainy', 'Overcast', 'Sunny', 'Sunny',
           'Rainy', 'Sunny', 'Overcast', 'Overcast', 'Rainy']
Temp = ['Hot', 'Hot', 'Hot', 'Mild', 'Cool', 'Cool', 'Cool', 'Mild', 'Cool', 'Mild', 'Mild', 'Mild', 'Hot', 'Mild']

Play = ['No', 'No', 'Yes', 'Yes', 'Yes', 'No', 'Yes', 'No', 'Yes', 'Yes', 'Yes', 'Yes', 'Yes', 'No']
```

```
# Import LabelEncoder
from sklearn import preprocessing

#creating labelEncoder
le = preprocessing.LabelEncoder()

# Converting string labels into numbers.
weather_encoded = le.fit_transform(Weather)
print(weather_encoded)
```

→ [2 2 0 1 1 1 0 2 2 1 2 0 0 1]

```
# Converting string labels into numbers
temp_encoded = le.fit_transform(Temp)
label = le.fit_transform(Play)
```

```
print("Temp:", temp_encoded)
print("Play :", label)
```

→ Temp: [1 1 1 2 0 0 0 2 0 2 2 2 1 2]
Play : [0 0 1 1 1 0 1 0 1 1 1 1 1 0]

```
# Combinig weather and temp into single listof tuples
features = [tup for tup in zip(weather_encoded, temp_encoded)]
print(features)
```

→ [(2, 1), (2, 1), (0, 1), (1, 2), (1, 0), (1, 0), (0, 0), (2, 2), (2, 0), (1, 2), (2, 2), (0, 2), (0, 1), (1, 2)]

```
#Import Gaussian Naïve Bayes model
from sklearn.naive_bayes import GaussianNB
```

```
#Create a Gaussian Classifier
model = GaussianNB()
```

```
# Train the model using the training sets
model.fit(features, label)
```

```
#Predict Output
predicted= model.predict([[0,2]]) # 0:Overcast, 2:Mild
print("Predicted Value:", predicted)
```

→ Predicted Value: [1]

✓ 2. In the following step, you will be working on building your own Naïve Bayes Classifier.

```
data = {
    'Tid': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
```

```

'Refund': ['Yes', 'No', 'No', 'Yes', 'No', 'No', 'Yes', 'No', 'No', 'No'],
'Marital Status': ['Single', 'Married', 'Single', 'Married', 'Divorced', 'Married', 'Divorced', 'Single', 'Married',
'Taxable Income': ['125K', '100K', '70K', '120K', '95K', '60K', '220K', '85K', '75K', '90K'],
'Evade': ['No', 'No', 'No', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'Yes']
}

df = pd.DataFrame(data)

df['Taxable Income'] = df['Taxable Income'].str.replace('K', '').astype('int64')
df['Income'] = pd.cut(df['Taxable Income'], bins=[0, 80, 150, np.inf], labels=['Low', 'Medium', 'High'])

df.drop('Taxable Income', axis=1, inplace=True)

```

```



from sklearn import preprocessing

le = preprocessing.LabelEncoder()

df['Refund'] = le.fit_transform(df['Refund'])
df['Marital Status'] = le.fit_transform(df['Marital Status'])
df['Income'] = le.fit_transform(df['Income'])
df['Evade'] = le.fit_transform(df['Evade'])

df

```

	Tid	Refund	Marital Status	Evade	Income	
0	1	1	2	0	2	
1	2	0	1	0	2	
2	3	0	2	0	1	
3	4	1	1	0	2	
4	5	0	0	1	2	
5	6	0	1	0	1	
6	7	1	0	1	0	
7	8	0	2	1	2	
8	9	0	1	0	1	
9	10	0	2	1	2	

Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

```

X = df[['Refund', 'Marital Status', 'Income']]
y = df['Evade']

```

```

class NaiveBayesClassifier:
    def fit(self, X, y):
        n_samples, n_features = X.shape
        print("No. of samples: ", n_samples, "\nNo. of features: ", n_features)
        self.classes = np.unique(y)
        n_classes = len(self.classes)
        print("No. of classes: ", n_classes)

        # Calculate priors
        self.priors = np.zeros(n_classes)
        for idx, c in enumerate(self.classes):
            self.priors[idx] = np.sum(y == c) / n_samples

        print("\nprior probability of NO: ", self.priors[0])
        print("prior probability of YES: ", self.priors[1])

        # Likelihoods with Laplace Smoothing
        self.likelihoods = {}
        for feature_idx in range(n_features):
            self.likelihoods[feature_idx] = {}
            for idx, c in enumerate(self.classes):
                X_c = X[y == c]

```

```

feature_values = np.unique(X[:, feature_idx])
self.likelihoods[feature_idx][c] = {}
for value in feature_values:
    self.likelihoods[feature_idx][c][value] = (np.sum(X_c[:, feature_idx] == value) + 1) / (len(X_c) + 1)


def predict(self, X):
    y_pred = []
    for x in X:
        posteriors = []
        for idx, c in enumerate(self.classes):
            prior = np.log(self.priors[idx])
            likelihood = np.sum([np.log(self.likelihoods[feature_idx][c].get(x[feature_idx], 1e-6)) for feature_idx in self.feature_idx])
            posterior = prior + likelihood
            posteriors.append(posterior)
        y_pred.append(self.classes[np.argmax(posteriors)])
    return np.array(y_pred)

```

```

classifier = NaiveBayesClassifier()
classifier.fit(X.values, y.values)

```

 No. of samples: 10
 No. of features: 3
 No. of classes: 2


prior probability of NO: 0.6
 prior probability of YES: 0.4



```

predictions = classifier.predict(X.values)
df['Predicted Evade'] = predictions

df[['Tid', 'Evade', 'Predicted Evade']]

```



	Tid	Evade	Predicted Evade	
0	1	0	0	
1	2	0	0	
2	3	0	0	
3	4	0	0	
4	5	1	1	
5	6	0	0	
6	7	1	1	
7	8	1	1	
8	9	0	0	
9	10	1	1	

```


X_test = np.array([
    [1, 0, 1], # X1: Refund = No, Status = Married, Income = Low (60K)
    [0, 1, 2] # X2: Refund = Yes, Status = Divorced, Income = Medium (90K)
])

```

```

y_test_pred = classifier.predict(X_test)
print("Class \nNO : 0, YES : 1")
print("Refund = No, Status = Married, Income = Low (60K), Predicted Class: ", y_test_pred[0])
print("Refund = Yes, Status = Divorced, Income = Medium (90K), Predicted Class: ", y_test_pred[1])

```

 Class
 NO : 0, YES : 1
 Refund = No, Status = Married, Income = Low (60K), Predicted Class: 0
 Refund = Yes, Status = Divorced, Income = Medium (90K), Predicted Class: 0